

# MACHINE LEARNING

Applied to French Football

*Predicting Ligue 1 Match Outcomes using AI*



---

Written by  
**Sami BENNANE**

Data & Artificial Intelligence Project

October 23, 2025

# 1 Introduction

## 1.1 Project Overview

This project explores the use of machine learning to predict the outcome of French Ligue 1 football matches. The problem is formulated as a **multi-class classification task** where the target variable  $y$  can take three possible values:

$$y \in \{\text{Home Win}, \text{Draw}, \text{Away Win}\}.$$

Given a set of features  $\mathbf{x} = (x_1, x_2, \dots, x_n)$  describing each match (teams, goals, etc.), the goal is to approximate a function  $f_\theta(\mathbf{x})$  parameterized by  $\theta$ :

$$\hat{y} = f_\theta(\mathbf{x}) \approx y.$$

### 1.1.1 Dataset Overview

The dataset  $\mathcal{D}$  contains historical results from the French Ligue 1, covering multiple seasons and totaling 7,378 matches. Each observation represents one match and includes both categorical and numerical data:

Home Team, Away Team, Home Team Goals, Away Team Goals, Winner.

Formally, the dataset can be written as:

$$\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^m, \quad m = 7378.$$

Here:

- $\mathbf{x}^{(i)} = (x_1^{(i)}, x_2^{(i)}, \dots, x_n^{(i)})$  is the **feature vector** describing match  $i$  (for example, the home team, away team, and goals scored),
- $y^{(i)}$  is the **true outcome** of match  $i$  (Home Win, Draw, or Away Win),
- $m$  is the **number of samples**, i.e. the total number of matches in the dataset.

In other words, the dataset  $\mathcal{D}$  contains  $m = 7378$  examples, each composed of a vector of  $n$  input features and a corresponding label  $y^{(i)}$ . Mathematically, this structure is represented as:

$$X \in \mathbb{R}^{m \times n}, \quad y \in \{0, 1, 2\}^m,$$

where each row of  $X$  corresponds to a specific Ligue 1 match and each entry in  $y$  encodes its outcome.

### 1.1.2 Feature Selection Rationale

In the original dataset, additional attributes were available such as:

- match date and stadium,
- referee name,
- competition round or season,

- audience or weather conditions.

However, many of these fields either:

1. lack a clear predictive relationship with match outcomes (e.g., referee name), or
2. introduce unnecessary noise or high-dimensional sparsity (hundreds of unique stadiums).

To keep the model simple and interpretable, we selected only the most informative and universally available features:

$\mathbf{X} = [\text{Home Team}, \text{Away Team}, \text{Home Team Goals}, \text{Away Team Goals}]$ .

This ensures the model focuses on fundamental competitive factors rather than context-dependent randomness.

### 1.1.3 Matrix Representation and Python Integration

After preprocessing, the data is represented in matrix form:

$$X \in \mathbb{R}^{m \times n}, \quad y \in \{0, 1, 2\}^m,$$

where:

- $m$  is the number of samples (i.e., the total number of matches, here  $m = 7378$ ),
- $n$  is the number of features describing each match (here  $n = 82$  after encoding),
- $X$  contains both numerical and one-hot encoded categorical variables,
- $y$  contains the encoded outcomes (Home Win = 0, Draw = 1, Away Win = 2).

Each row of  $X$  corresponds to one Ligue 1 match, and each column corresponds to a measurable or categorical characteristic of that match. For example:

$$X^{(i)} = [\text{Home\_Lyon}, \text{Away\_PSG}, \text{Goals\_Home} = 1, \text{Goals\_Away} = 2, \dots]$$

represents all the information associated with match  $i$  between Lyon and Paris SG.

In Python, this transformation is achieved using `pandas.get_dummies()` for categorical variables and `LabelEncoder()` for the target variable:

$$X_{\text{final}} = [\text{One-Hot}(\text{Home Team}), \text{One-Hot}(\text{Away Team}), \text{Goals\_Home}, \text{Goals\_Away}]$$

This results in a rectangular matrix  $X$  of shape  $(m, n)$  ready to be used in vectorized linear algebra operations.

This matrix representation is essential since most machine learning algorithms—including logistic regression—are defined as vectorized models:

$$\hat{\mathbf{y}} = \sigma(X\boldsymbol{\theta}),$$

where:

- $X$  is the **design matrix** (the input data),
- $\boldsymbol{\theta}$  is the **parameter vector** learned during training,
- $\sigma$  is the **activation function** (softmax in the multiclass case).

This explicit matrix formulation allows efficient computation with NumPy and `scikit-learn`, bridging the theoretical foundation of machine learning with its practical implementation in Python.

### 1.1.4 Football Interpretation

From a football perspective:

- The team identifiers in  $X$  encode each club's identity and past performance trends.
- The goal statistics provide a quantitative link between offensive/defensive strength and the match outcome.
- The model captures realistic patterns, such as the **home advantage effect** or the dominance of top teams like PSG.

By combining mathematical modeling and domain intuition, the project aims to connect statistical learning with real-world football dynamics.

## 2 Mathematical Model

### 2.1 Logistic Regression: Theory

For binary classification, logistic regression models the probability:

$$P(y = 1 \mid \mathbf{x}) = \sigma(\mathbf{x}^\top \boldsymbol{\theta}) = \frac{1}{1 + e^{-\mathbf{x}^\top \boldsymbol{\theta}}}.$$

For  $K$  outcomes (Home, Draw, Away), the model generalizes via the **softmax function**:

$$P(y = k \mid \mathbf{x}) = \frac{e^{\mathbf{x}^\top \boldsymbol{\theta}_k}}{\sum_{j=1}^K e^{\mathbf{x}^\top \boldsymbol{\theta}_j}}, \quad k = 1, \dots, K,$$

where each  $\boldsymbol{\theta}_k$  corresponds to a class (e.g., Home Win).

## 3 Matrix Formulation and Python Implementation

### 3.1 Matrix Representation

Let:

$$X = \begin{bmatrix} | & | & & | \\ \mathbf{x}^{(1)} & \mathbf{x}^{(2)} & \dots & \mathbf{x}^{(m)} \\ | & | & & | \end{bmatrix}^\top \in \mathbb{R}^{m \times n}, \quad \boldsymbol{\Theta} = [\boldsymbol{\theta}_1 \quad \boldsymbol{\theta}_2 \quad \dots \quad \boldsymbol{\theta}_K] \in \mathbb{R}^{n \times K}.$$

The model outputs:

$$\hat{Y} = \text{softmax}(X\boldsymbol{\Theta}) \in \mathbb{R}^{m \times K},$$

where each row  $\hat{y}^{(i)}$  is a probability vector over all  $K$  classes.

The vectorized **cost function** is:

$$J(\boldsymbol{\Theta}) = -\frac{1}{m} \text{Tr} \left( Y^\top \log(\hat{Y}) \right),$$

and its gradient:

$$\nabla_{\boldsymbol{\Theta}} J = \frac{1}{m} X^\top (\hat{Y} - Y).$$

These compact equations show how all  $m$  examples are processed in parallel through efficient matrix operations.

## 3.2 Python Implementation

Listing 1: Vectorized implementation of multinomial logistic regression.

```
# X: input matrix (m, n)
# Theta: parameters (n, K)
# y_onehot: true labels in one-hot encoding (m, K)

z = X @ Theta # Linear step
y_hat = np.exp(z) / np.sum(np.exp(z), axis=1, keepdims=True) # Softmax
cost = -(1/m) * np.sum(y_onehot * np.log(y_hat)) # Loss

grad = (1/m) * X.T @ (y_hat - y_onehot) # Gradient
Theta -= learning_rate * grad # Update rule
```

This implementation corresponds exactly to the analytical formula:

$$\Theta := \Theta - \eta \nabla_{\Theta} J(\Theta),$$

where  $\eta$  is the learning rate.

## 4 Data Processing

Categorical variables such as *Home Team* and *Away Team* were encoded via **one-hot encoding**:

$$X_{\text{encoded}} = [\text{Home\_Team\_Lyon}, \text{Away\_Team\_PSG}, \dots, \text{Goals\_Home}, \text{Goals\_Away}].$$

After preprocessing:

$$X \in \mathbb{R}^{7378 \times 82}, \quad y \in \{0, 1, 2, \dots, 40\}.$$

The data was then split into 80% training and 20% testing subsets.

## 5 Results and Analysis

The multinomial logistic regression achieved a test accuracy of:

$$\text{Accuracy} = 54.47\%.$$

This performance, although modest, reflects the inherent unpredictability of football results. Nevertheless, the model successfully captured strong tendencies such as *home advantage* and the dominance of major teams.

### 5.1 Interpretation

The model parameters  $\theta_{\text{home}}$  and  $\theta_{\text{away}}$  act as learned weights reflecting each team's influence. The predicted class is given by:

$$\hat{y} = \arg \max_k \mathbf{x}^\top \theta_k.$$

Teams with larger positive weights (e.g., PSG, Lyon) tend to yield higher predicted probabilities for winning at home.

## 6 Evaluation Metrics and Football Interpretation

Evaluating a machine learning model is crucial to understanding both its predictive power and its limitations. In this project, we focus on the most relevant metrics for football match prediction: **Accuracy**, **Precision**, **Recall**, **F1-score**, and their aggregated versions (**macro** and **weighted averages**).

### 6.1 Accuracy

The simplest and most intuitive metric, accuracy measures the overall proportion of correct predictions:

$$\text{Accuracy} = \frac{\text{Number of Correct Predictions}}{\text{Total Number of Predictions}}.$$

For instance, if the model correctly predicts 800 matches out of 1476:

$$\text{Accuracy} = \frac{800}{1476} \approx 54.47\%.$$

In football terms, this means the model “guesses” the right outcome slightly more than once every two matches.

### 6.2 Precision and Recall

To evaluate the quality of predictions for each possible outcome (Home Win, Draw, Away Win), we use:

$$\text{Precision} = \frac{TP}{TP + FP}, \quad \text{Recall} = \frac{TP}{TP + FN}.$$

- **Precision** indicates how often the model is correct when it predicts a given result. For example, when the model predicts a *home win*, how often is it right?
- **Recall** measures how well the model identifies all true occurrences. For example, among all matches that actually ended with a home win, how many did the model correctly detect?

The harmonic mean of both gives the **F1-score**:

$$F1 = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}.$$

### 6.3 Macro and Weighted Averages

Because some classes (teams or results) appear more often than others, we use two types of averages:

$$\begin{aligned} \text{Macro Average} &= \frac{1}{K} \sum_{k=1}^K \text{Metric}_k \\ \text{Weighted Average} &= \frac{\sum_{k=1}^K n_k \times \text{Metric}_k}{\sum_{k=1}^K n_k} \end{aligned}$$

where  $n_k$  is the number of samples (matches) in class  $k$ .

- **Macro average** gives equal importance to all classes — useful to see if the model treats small clubs (e.g., Metz, Brest) fairly.
- **Weighted average** gives more weight to frequent classes — which in football means the model performance is driven by the big clubs that dominate the dataset (e.g., PSG, Lyon, Marseille).

## 6.4 Football Interpretation

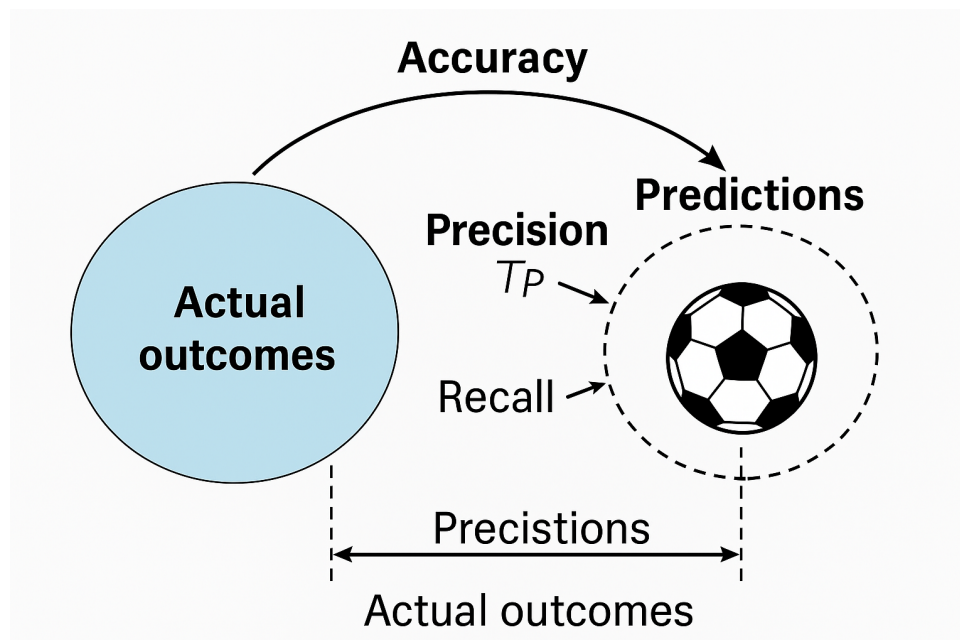
In our Ligue 1 dataset:

- The **macro average** is slightly lower than the accuracy, showing that smaller or less frequent outcomes (like draws or away wins) are harder to predict.
- The **weighted average** is close to the accuracy, indicating that the model is consistent on teams with many games (the major clubs).

This means the model captures strong patterns such as:

- Home advantage (more likely to predict a home win correctly).
- The dominance of top teams (PSG, Lyon, Marseille are predicted more accurately).
- Greater uncertainty for balanced or rare outcomes (draws, small teams winning away).

## 6.5 Visual Representation



**Figure:** Conceptual diagram showing the relationship between predictions and actual outcomes. A perfect model would have precision and recall equal to 1. In football prediction, some imprecision is inevitable due to random events (injuries, referee decisions, etc.).

## 7 Conclusion

This project demonstrates the full pipeline of a machine learning approach — from data preprocessing to mathematical modeling and numerical optimization.

Key takeaways:

- The mathematical model can be entirely expressed in **matrix form**.
- NumPy seamlessly implements these vectorized operations, achieving high computational efficiency.
- Extending this model to neural networks would only add more layers of similar matrix multiplications.

Future improvements could include:

- Adding regularization:

$$J_{\text{reg}}(\Theta) = J(\Theta) + \lambda \|\Theta\|_2^2$$

- Using ensemble models or neural networks.
- Integrating temporal dynamics of teams with LSTM-based models.

**Keywords:** Machine Learning, Logistic Regression, Matrix Formulation, NumPy, Sports Analytics, French Ligue 1



## Appendix — Complete Gradient Derivation via Differentials

We now derive the gradient of the multiclass logistic regression cost function using the **differential approach**, providing a step-by-step derivation with all details.

### 1. Problem Setup and Definitions

Let:

$$J(W) = \frac{1}{m} \sum_{i=1}^m L_i, \quad L_i = - \sum_{k=1}^K Y_{ik} \log(\hat{Y}_{ik}), \quad \hat{Y} = \text{softmax}(Z), \quad Z = XW$$

where:

- $X \in \mathbb{R}^{m \times n}$  is the data matrix (m examples, n features)
- $W \in \mathbb{R}^{n \times K}$  are the model parameters
- $Y, \hat{Y} \in \mathbb{R}^{m \times K}$  are the true and predicted label matrices
- $Z \in \mathbb{R}^{m \times K}$  are the logits

### 2. Step 1: Gradient of Loss with Respect to Logits

First, we compute  $\frac{\partial J}{\partial Z}$ .

#### 2.1 For one example $L_i$ :

$$\frac{\partial L_i}{\partial z_j} = - \sum_{k=1}^K Y_{ik} \frac{1}{\hat{Y}_{ik}} \frac{\partial \hat{Y}_{ik}}{\partial z_j}$$

#### 2.2 Using softmax derivatives:

$$\frac{\partial \hat{Y}_{ik}}{\partial z_j} = \hat{Y}_{ik} (\delta_{kj} - \hat{Y}_{ij})$$

where  $\delta_{kj}$  is the Kronecker delta.

**2.3 Substitution and simplification:**

$$\begin{aligned}
\frac{\partial L_i}{\partial z_j} &= - \sum_{k=1}^K Y_{ik} \frac{1}{\hat{Y}_{ik}} \hat{Y}_{ik} (\delta_{kj} - \hat{Y}_{ij}) \\
&= - \sum_{k=1}^K Y_{ik} (\delta_{kj} - \hat{Y}_{ij}) \\
&= -Y_{ij} + \hat{Y}_{ij} \sum_{k=1}^K Y_{ik} \\
&= \hat{Y}_{ij} - Y_{ij} \quad (\text{since } \sum_{k=1}^K Y_{ik} = 1 \text{ for one-hot vectors})
\end{aligned}$$

**2.4 For all examples (matrix form):**

$$\frac{\partial J}{\partial Z} = \frac{1}{m}(\hat{Y} - Y)$$

**3. Step 2: Differential of the Cost Function**

Using the definition of differential via inner product:

$$dJ = \left\langle \frac{\partial J}{\partial Z}, dZ \right\rangle = \text{Tr} \left( \left( \frac{\partial J}{\partial Z} \right)^\top dZ \right)$$

Substitute  $\frac{\partial J}{\partial Z} = \frac{1}{m}(\hat{Y} - Y)$ :

$$dJ = \left\langle \frac{1}{m}(\hat{Y} - Y), dZ \right\rangle$$

**4. Step 3: Relating  $dZ$  and  $dW$** 

Since  $Z = XW$  and  $X$  is constant:

$$dZ = X \cdot dW$$

Substitute into the differential:

$$dJ = \left\langle \frac{1}{m}(\hat{Y} - Y), X \cdot dW \right\rangle$$

**5. Step 4: Matrix Inner Product Manipulation**

We use the property of matrix inner products:

$$\langle A, BC \rangle = \langle B^\top A, C \rangle$$

Apply this with  $A = \frac{1}{m}(\hat{Y} - Y)$ ,  $B = X$ ,  $C = dW$ :

$$dJ = \left\langle X^\top \cdot \frac{1}{m}(\hat{Y} - Y), dW \right\rangle$$

## 6. Step 5: Identification of the Gradient

By definition of the gradient in matrix calculus:

$$dJ = \langle \nabla_W J, dW \rangle$$

Comparing with our expression:

$$\langle \nabla_W J, dW \rangle = \left\langle \frac{1}{m} X^\top (\hat{Y} - Y), dW \right\rangle$$

Since this holds for all  $dW$ , we can identify:

$$\boxed{\nabla_W J = \frac{1}{m} X^\top (\hat{Y} - Y)}$$

## 7. Dimension Verification

- $X^\top \in \mathbb{R}^{n \times m}$
- $(\hat{Y} - Y) \in \mathbb{R}^{m \times K}$
- $X^\top (\hat{Y} - Y) \in \mathbb{R}^{n \times K}$  (matches dimensions of  $W$ )

## 8. Interpretation and Usage

The differential approach shows that:

$$dJ = \langle \nabla_W J, dW \rangle$$

meaning  $\nabla_W J$  gives the direction of steepest ascent of  $J$ .

In gradient descent:

$$W \leftarrow W - \eta \nabla_W J$$

parameters are updated opposite to the gradient, reducing the loss.

### Key Insights:

- The Kronecker delta property and one-hot encoding simplification are crucial
- Matrix inner products provide elegant notation for chain rule applications
- The final result  $\nabla_W J = \frac{1}{m} X^\top (\hat{Y} - Y)$  is remarkably simple
- This approach generalizes naturally to neural networks and backpropagation